$$e = \mathrm{mean}_{mask,T}(H), \quad r = \mathrm{LN}(\mathrm{GELU}(eW_{pre}^{\mathsf{T}})), \quad z = rW_{cls}^{\mathsf{T}} + b$$

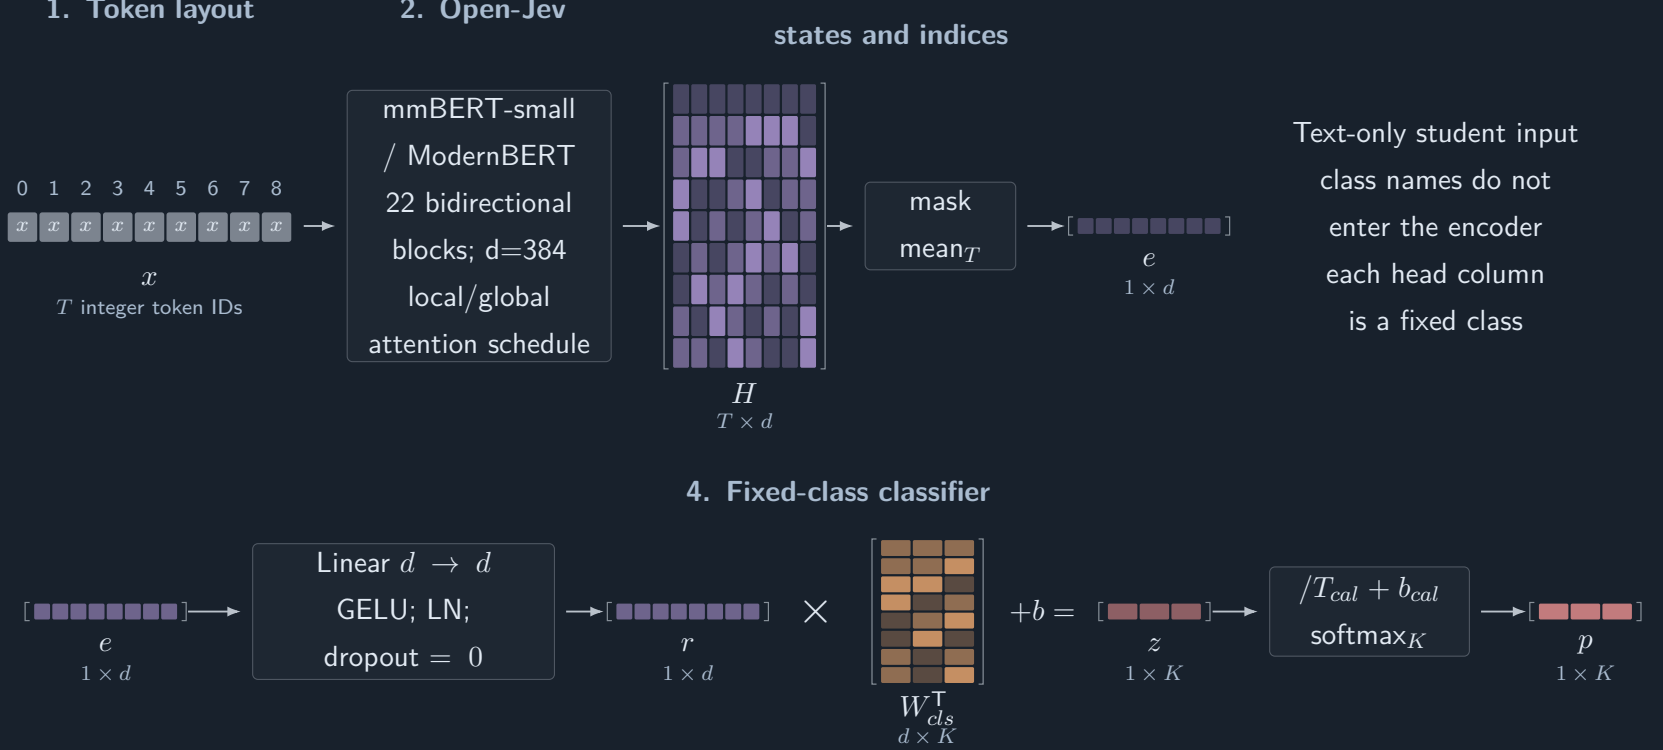

Axes $d = 384$, K fixed; Single-question illustration; batch omitted. Tile counts encode axes, not literal sizes.

Objects Gray: input/context; violet: hidden states; orange: parameter matrices; red: logits/probabilities. Positions, code IDs and owner maps are integers, not scores. Tiles show roles, not measured magnitudes.

Mechanism Masked mean, Linear-GELU-LayerNorm and a fixed K-class head. Runtime option text is absent; new class names require a new task mapping/head, not just a prompt edit.

Default student minimizes KL from teacher probabilities; optional gold CE has weight zero by default. Internal classifier is a pinned Transformers structural reference.