$$\widehat{H} = H[P,:], \quad U = \mathrm{GELU}(\widehat{H}W_1^{\mathsf{T}} + b_1), \quad s = \mathrm{squeeze}(UW_2^{\mathsf{T}} + b_2)$$

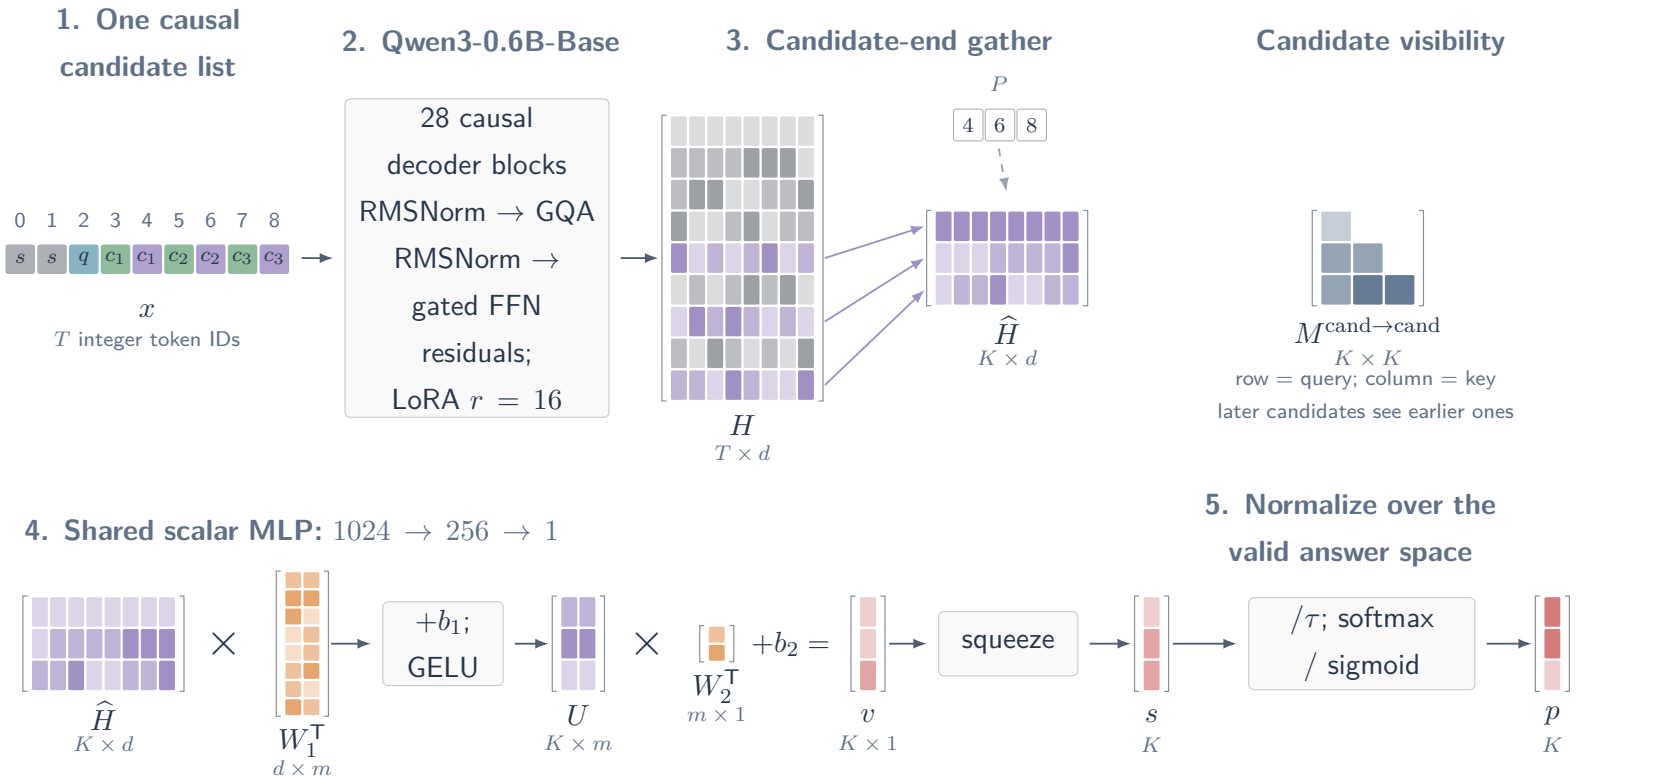

Axes	One question shown; batch B omitted. T: tokens; K: valid candidates; $d = 1024$, $m = 256$. Tiles illustrate axis structure, not the literal dimensions.
Objects	x and P are integer IDs/positions; $H, \widehat{H}, U$ are continuous states; W_1, W_2 are learned weights; s are raw scores and p probabilities. Biases appear in the formula.
Mechanism	Choice example, `candidate_scalar/list`: actual candidate-end tokens, no inserted markers. Noul instead selects the final Answer cue ($K = 1$). Choice/score use masked softmax; noul/multi-label use sigmoid. No tree cache; LM head bypassed.