$$W_{slot}^{(0)} = W_{lm}[\ell, :], \qquad z = hW_{slot}^{\mathsf{T}}, \qquad p = \mathrm{softmax}_k(z/\tau)$$

Initialize a separate slot head from selected language-model head rows

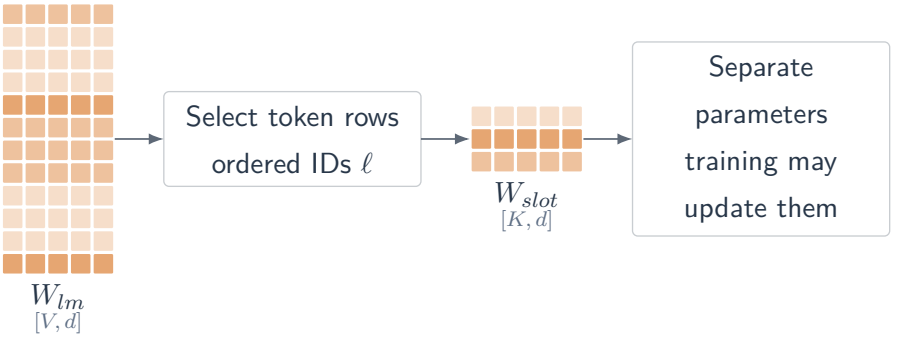

Project into a fixed set of output slots

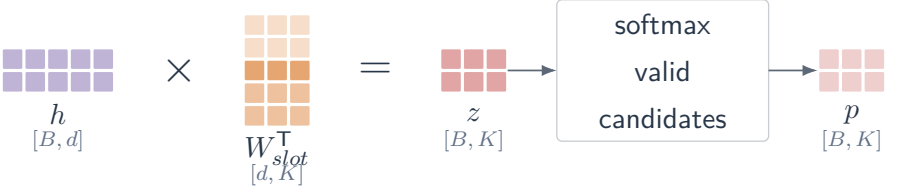

Axes $W_{lm} : [V, d]$; $W_{slot} : [K, d]$; $h : [B, d]$. K is the configured slot capacity. Initialization copies content, not a live link to the original rows.

Meaning This comparison family changes the output parameterization. Attention visibility is a separate design choice: causal or bidirectional variants must be checked per implementation. A standard causal decoder is not automatically equivalent to a modified mask.

Cell counts illustrate axes, not model sizes; tint identifies roles, not numeric values.