$$r_{1:G} \sim \pi_\theta(\cdot \mid x), \quad X' = [x; r; \text{answer cue}], \quad p = \text{softmax}_k(z_{\ell_k}/\tau)$$

Generate rationale tokens with autoregressive dependencies

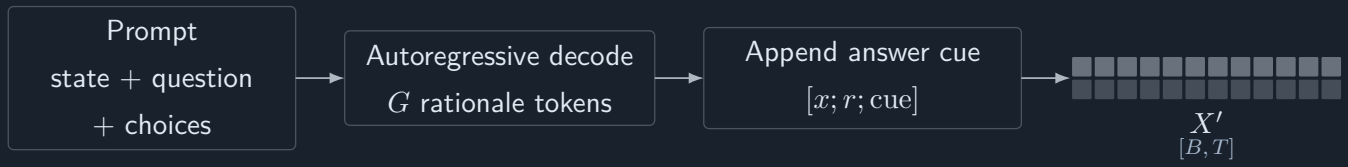

Read answer-token probabilities after the generated rationale

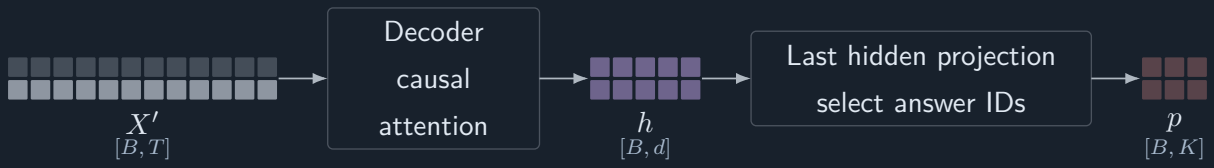

Axes	G: generated rationale length; T: total length after concatenation; K: answer choices. The final vocabulary selection has the same shape as answer-token readout.
Meaning	This is a comparison, not an implemented MetisCraft fallback. Rationale tokens add serial decoding dependencies; cache reuse, batching and the final readout determine actual model invocations. Sampling can change the resulting distribution.

Cell counts illustrate axes, not model sizes; tint identifies roles, not numeric values.