$$\text{input layout} \quad + \quad \text{information flow} \quad + \quad \text{score operator} \quad \longrightarrow \quad \text{typed judgment}$$

Output parameterization: vocabulary tokens, fixed slots or candidate scalar scores

1. Answer-token logits
$[B, V] \ \rightarrow \ [B, L]$
gather selected vocabulary IDs

2. Fixed slot head
$[B, d]\,[d, K] \ \rightarrow \ [B, K]$
copy token rows at initialization

3. Candidate scalar head
$[R, d] \ \rightarrow \ [R] \ \rightarrow \ [B, K]$
independent pair rows

Structured or encoder alternatives: score operators and information flow differ

4. Pointer head
$q \ : \ [B, d_p], \ k \ : \ [B, K, d_p]$
decision–option dot products

5. Joint schema head
$O : [C, Q, K, w], \ F : [C, Q, w]$
evidence routing
+ field interaction

6. Encoder scoring
uni / bi / cross layouts
visibility depends on the layout

Execution and scope: distinguish a comparison family from a shipped implementation

7. Generate then read
rationale + answer-token readout
serial decoding dependencies

MetisCraft scalar default: list
$[B, K, d] \ \rightarrow \ [B, K]$
shared causal sequence

Implemented: answer tokens + scalar
tree packing: not implemented
no universal speed ranking

Axes B: questions; $R = \sum_b K_b$: independent candidate rows; C: packed sequences; Q: questions per sequence; K: candidate capacity; L: answer-token count. Dots indicate operations, not benchmark results.

Meaning The seven families compare mechanisms; they are not seven supported MetisCraft backends. The shipped toolkit implements answer-token and candidate-scalar readouts, with list and pairs layouts for the latter. Accuracy and latency require matched measurements.

Cell counts illustrate axes, not model sizes; tint identifies roles, not numeric values.