$$h_{b,k} = H_{b,P_{b,k},:}, \quad u = \mathrm{GELU}(hW_1^\mathsf{T} + b_1), \quad S = \mathrm{squeeze}(uW_2^\mathsf{T} + b_2)$$

Default scalar layout: all candidates share one causal sequence

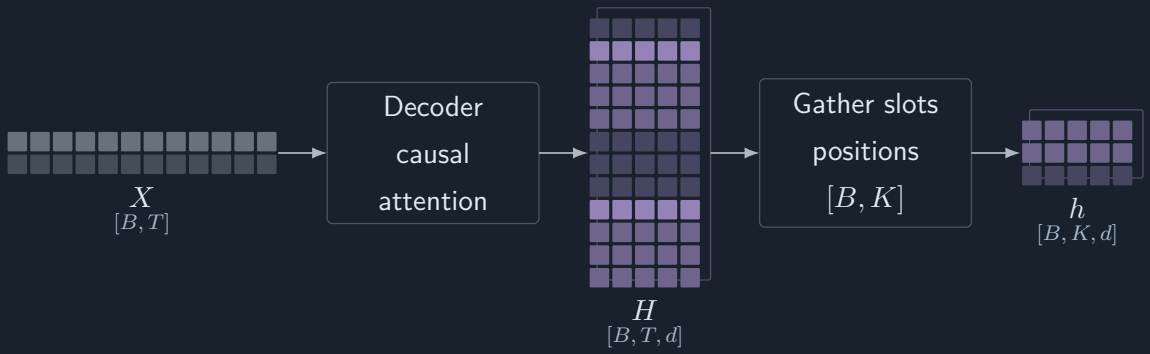

Gather candidate slots; apply the same scalar head at every slot

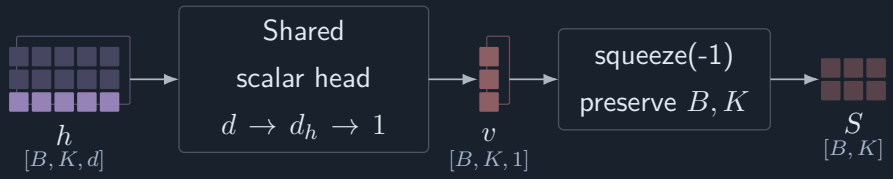

Axes $P_{b,k} \in \{0, \ldots, T-1\}$ selects a candidate's slot after left padding. $h_{b,k} = H_{b,P_{b,k},:}$ has shape $[d]$.

Meaning Later candidates can attend to earlier candidates under the causal mask. An independent sigmoid output does not make their hidden representations independent. List is the current default; tree packing is not implemented.

Cell counts illustrate axes, not model sizes; tint identifies roles, not numeric values.